

# The Blind Spot Paradox


When a self-repairing model destroys the evidence its own monitor needs

---

Alexandre Boyer   Mélaine Gouillou   Salomé Fonvielle   Ulysse Petit-Tichanné



Research Track, Topic 39, CentraleSupélec. Supervisor: Raphaël Minato

# Concept drift: the rule changes, the data does not



A model trained on the left keeps predicting, confidently, and is now **wrong on the shaded band**.

# Defence one: a model that repairs itself




This is the **Adaptive Random Forest**. Ten trees, replaced one at a time.

## Defence two: a watchdog that raises the alarm


the model answers, one observation at a time (x = wrong)

x o o o x o o o o x o o o x o o x x o x x o x




Two communities, two tools, never measured against each other.

# Put them together, and they cancel out



The repair has a date. The alarm never gets one. (schematic)


## Why not just lower the threshold?



A low threshold also fires when **nothing has happened**, and an alarm that cries wolf gets switched off.

# Our contribution: when is the forest actually repaired?

The alarm is what we *want*. This date measures what **stops it** from ever firing.



It fires at the *first* tree. It times the start of a repair, not its end.

## Where we stand

- Yes** the alarm really does stay silent, and we measured how often. Turning the threshold down does not rescue it, it only trades silence for false alarms.
- But** the stopwatch the field uses starts at the **first tree out of ten**. It makes the forest look faster than it is, so every published figure about this race is optimistic.
- Next** put the model and the watchdog on one shared clock, and say which setting a practitioner should actually pick.

Thank you.